

```

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import plotly.express as px
from plotly.offline import iplot, plot
from plotly.subplots import make_subplots
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.metrics import ConfusionMatrixDisplay
from sklearn.preprocessing import MinMaxScaler
from sklearn.svm import SVC
from sklearn.ensemble import RandomForestClassifier
from sklearn.linear_model import LogisticRegression
import warnings
warnings.filterwarnings('ignore')

```

```
df = pd.read_csv("/content/train.csv")
```

```
df.sample(5)
```

```
{"type": "dataframe"}
```

```
print(f"Number of Row : {df.shape[0]}\nNumber of Columns : {df.shape[1]}")
```

```
Number of Row : 2000
```

```
Number of Columns : 21
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 2000 entries, 0 to 1999
```

```
Data columns (total 21 columns):
```

#	Column	Non-Null Count	Dtype
0	battery_power	2000 non-null	int64
1	blue	2000 non-null	int64
2	clock_speed	2000 non-null	float64
3	dual_sim	2000 non-null	int64
4	fc	2000 non-null	int64
5	four_g	2000 non-null	int64
6	int_memory	2000 non-null	int64
7	m_dep	2000 non-null	float64
8	mobile_wt	2000 non-null	int64
9	n_cores	2000 non-null	int64
10	pc	2000 non-null	int64
11	px_height	2000 non-null	int64
12	px_width	2000 non-null	int64
13	ram	2000 non-null	int64
14	sc_h	2000 non-null	int64
15	sc_w	2000 non-null	int64

16	talk_time	2000	non-null	int64
17	three_g	2000	non-null	int64
18	touch_screen	2000	non-null	int64
19	wifi	2000	non-null	int64
20	price_range	2000	non-null	int64

dtypes: float64(2), int64(19)

memory usage: 328.2 KB

```
df.isna().sum()
```

battery_power	0
blue	0
clock_speed	0
dual_sim	0
fc	0
four_g	0
int_memory	0
m_dep	0
mobile_wt	0
n_cores	0
pc	0
px_height	0
px_width	0
ram	0
sc_h	0
sc_w	0
talk_time	0
three_g	0
touch_screen	0
wifi	0
price_range	0

dtype: int64

```
df.describe()
```

```
{"type": "dataframe"}
```

```
pd.DataFrame({'Count': df.shape[0],
              'Null': df.isnull().sum(),
              'Null %': df.isnull().mean() * 100,
              'Cardinality': df.nunique()
            })
```

```
{"summary": "{\n  \"name\": \"\"\n}\",\n  \"rows\": 21,\n  \"fields\": [\n    {\n      \"column\": \"Count\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 0,\n        \"min\": 2000,\n        \"max\": 2000,\n        \"num_unique_values\": 1,\n        \"samples\": [\n          2000\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    },\n    {\n      \"column\": \"Null\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 0,\n        \"min\": 0,\n        \"max\": 0,\n        \"num_unique_values\": 1,\n        \"samples\": [\n          0\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    }\n  ]\n}
```

```

\"max\": 0,\n          \"num_unique_values\": 1,\n          \"samples\": [\n            0\n          ],\n          \"semantic_type\": \"\",\n          \"description\": \"\",\n          \"column\":\n          {\n            \"dtype\": \"number\",\n            \"std\": 0.0,\n            \"min\": 0.0,\n            \"max\": 0.0,\n            \"num_unique_values\": 1,\n            \"samples\": [\n              0.0\n            ],\n            \"semantic_type\": \"\",\n            \"description\": \"\",\n            \"column\": \"Cardinality\",\n            \"dtype\": \"number\",\n            \"std\": 493,\n            \"min\": 2,\n            \"max\": 1562,\n            \"num_unique_values\": 15,\n            \"samples\": [\n              1137\n            ],\n            \"semantic_type\": \"\",\n            \"description\": \"\"\n          }\n        ],\n      },\n    ],\n  },\n  \"type\": \"dataframe\"}

```

```
df.duplicated().any()
```

```
False
```

```
df_battery_price = df.groupby('price_range')['battery_power'].mean()
```

```

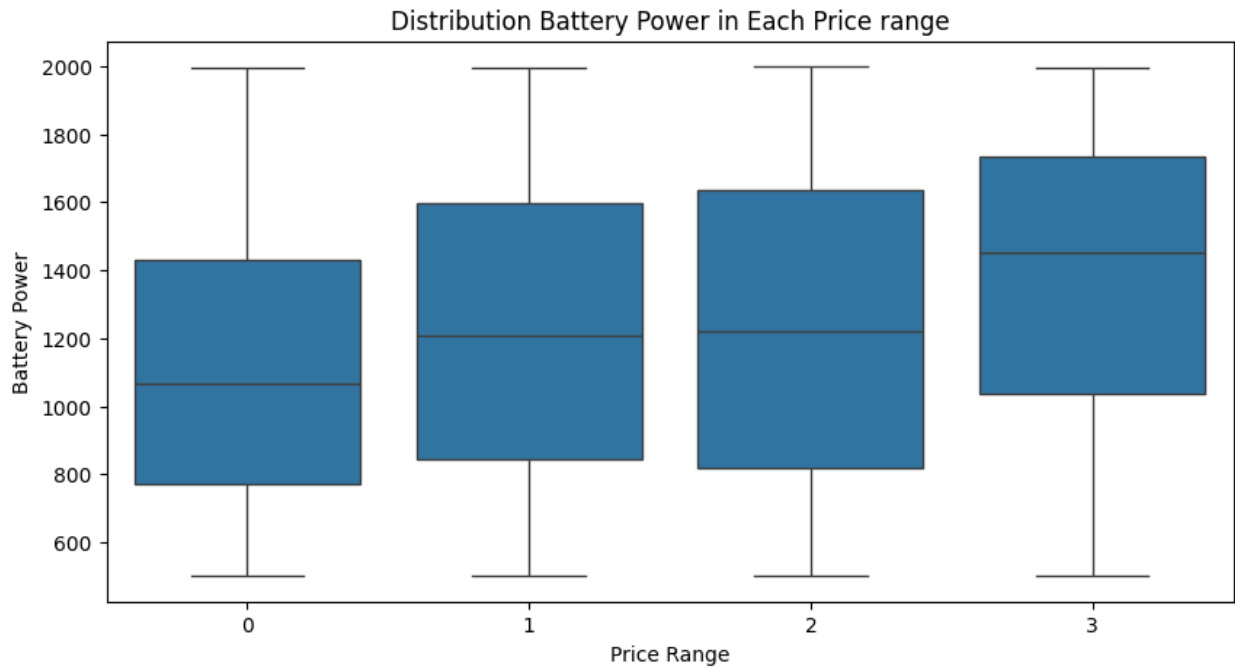
iplot(px.line(df_battery_price,
              labels={'value': 'Mean of Battery
Capacity', 'price_range': 'Price Range'},
              color_discrete_sequence=['red']
              ))

```

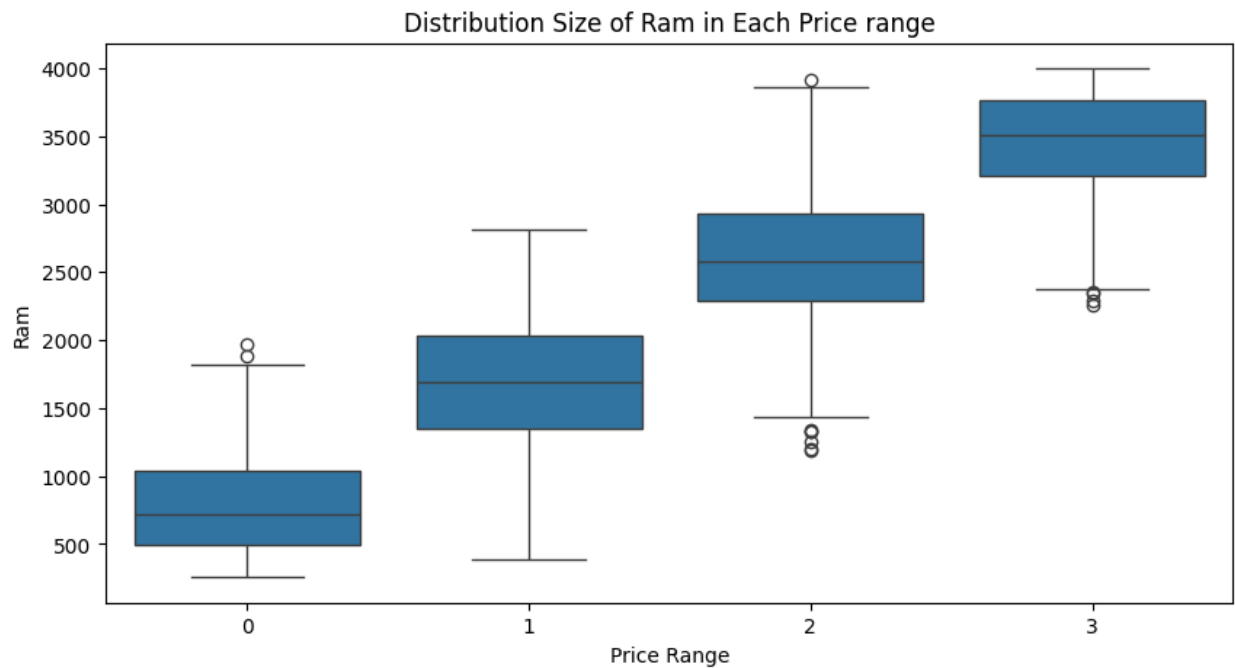
```

plt.figure(figsize=(10,5))
plt.title('Distribution Battery Power in Each Price range')
sns.boxplot(x=df['price_range'],y=df['battery_power'])
plt.xlabel('Price Range')
plt.ylabel('Battery Power')
plt.show()

```



```
plt.figure(figsize=(10,5))
plt.title('Distribution Size of Ram in Each Price range')
sns.boxplot(x=df['price_range'],y=df['ram'])
plt.xlabel('Price Range')
plt.ylabel('Ram')
plt.show()
```



```

df_4g = df['four_g'].value_counts()
iplob(px.pie(values=df_4g,
              names=['Support 4G', 'Not Support 4G'],
              template='plotly_dark',
              title='Is Support 4G ?'
              ).update_traces(textinfo='label+percent'))

df_3g = df['three_g'].value_counts()
iplob(px.pie(values=df_3g,
              names=['Support 3G', 'Not Support 3G'],
              template='plotly_dark',
              title='Is Support 3G ?'
              ).update_traces(textinfo='label+percent'))

x = df.drop(columns='price_range')
y = df.price_range

scaler = MinMaxScaler()
x = scaler.fit_transform(x)

x_train , x_test , y_train , y_test =
train_test_split(x,y,test_size=0.2)

print(f'Shape of X_Train {x_train.shape}')
print(f'Shape of X_Test {x_test.shape}')
print(f'Shape of Y_Train {y_train.shape}')
print(f'Shape of Y_Test {y_test.shape}')

Shape of X_Train (1600, 20)
Shape of X_Test (400, 20)
Shape of Y_Train (1600,)
Shape of Y_Test (400,)

model_params = {
    'svm':{
        'model' : SVC(gamma='auto'),
        'params':{
            'C':[1,10,20],
            'kernel':['rbf', 'linear']
        }
    },
    'random_forest':{
        'model':RandomForestClassifier(),
        'params':{
            'n_estimators':[1,5,10]
        }
    },
    'logistic_regression':{
        'model':LogisticRegression(solver='liblinear',multi_class='auto'),
        'params':{

```

```

        'C':[1,5,10]
    }
}

scores = []

for model_name , mp in model_params.items():
    clf =
GridSearchCV(mp[ 'model' ],mp[ 'params' ],cv=5,return_train_score=False)
    clf.fit(x,y)
    scores.append({
        'model':model_name,
        'best_scores':clf.best_score_,
        'best_params':clf.best_params_
    })

pd.DataFrame(scores,columns=[ 'model' , 'best_scores' , 'best_params' ])

{"summary":{"\n  \"name\": \"pd\", \n  \"rows\": 3, \n  \"fields\": [\n    {\n      \"column\": \"model\", \n      \"properties\": {\n        \"dtype\": \"string\", \n        \"num_unique_values\": 3, \n        \"samples\": [\n          \"svm\", \n          \"random_forest\", \n          \"logistic_regression\" \n        ], \n        \"semantic_type\": \"\", \n        \"description\": \"\" \n      }, \n      {\n        \"column\": \"best_scores\", \n        \"properties\": {\n          \"dtype\": \"number\", \n          \"std\": 0.08927672335683767, \n          \"min\": 0.7965, \n          \"max\": 0.9674999999999999, \n          \"num_unique_values\": 3, \n          \"samples\": [\n            0.9674999999999999, \n            0.7965, \n            0.8375 \n          ], \n          \"semantic_type\": \"\", \n          \"description\": \"\" \n        }, \n        {\n          \"column\": \"best_params\", \n          \"properties\": {\n            \"dtype\": \"object\", \n            \"semantic_type\": \"\", \n            \"description\": \"\" \n          } \n        } \n      ] \n    } \n  ], \"type\": \"dataframe\"}

model_svm = SVC(kernel='linear',C=20)
model_svm.fit(x_train,y_train)

SVC(C=20, kernel='linear')

score_svm_train = model_svm.score(x_train,y_train)
print(f"Train accuracy: {score_svm_train}")


Train accuracy: 0.9825

score_svm_test = model_svm.score(x_test,y_test)
print(f"Test accuracy: {score_svm_test}")

Test accuracy: 0.95

```

```
ConfusionMatrixDisplay.from_estimator(model_svm,
                                     x_test,
                                     y_test);
```



```
model_LR = LogisticRegression(C=10)
model_LR.fit(x_train,y_train)

LogisticRegression(C=10)


score_LR_train = model_LR.score(x_train,y_train)
print(f"Train accuracy: {score_LR_train}")

Train accuracy: 0.979375

score_LR_test = model_LR.score(x_test,y_test)
print(f"Test accuracy: {score_LR_test}")

Test accuracy: 0.935

ConfusionMatrixDisplay.from_estimator(model_LR,
                                     x_test,
                                     y_test);
```



```
model_RFC = RandomForestClassifier(n_estimators=10,random_state=42)
model_RFC.fit(x_train,y_train)
```

```
RandomForestClassifier(n_estimators=10, random_state=42)
```

```
score_RFC_train = model_RFC.score(x_train,y_train)
print(f"Train accuracy: {score_RFC_train}")
```

```
Train accuracy: 0.996875
```

```
score_RFC_test = model_RFC.score(x_test,y_test)
print(f"Test accuracy: {score_RFC_test}")
```

```
Test accuracy: 0.8125
```

```
from sklearn.model_selection import train_test_split
from sklearn.ensemble import GradientBoostingRegressor
import xgboost as xgb
from sklearn.metrics import mean_absolute_error, mean_squared_error,
r2_score
```

```
# Training Gradient Boosting model
```

```
gb_model = GradientBoostingRegressor(n_estimators=100,
learning_rate=0.1, max_depth=3, random_state=42)
```



```

gb_model.fit(x_train, y_train)

# Training XGBoost model
xgb_model = xgb.XGBRegressor(n_estimators=100, learning_rate=0.1,
max_depth=3, random_state=42)
xgb_model.fit(x_train, y_train)

# Making predictions on the testing set
gb_predictions = gb_model.predict(x_test)
xgb_predictions = xgb_model.predict(x_test)

mae = mean_absolute_error(y_test, gb_predictions)
print("Mean Absolute Error (MAE):", mae)

# Calculate Mean Squared Error (MSE)
mse = mean_squared_error(y_test, gb_predictions)
print("Mean Squared Error (MSE):", mse)

# Calculate Root Mean Squared Error (RMSE)
rmse = mean_squared_error(y_test, gb_predictions, squared=False)
print("Root Mean Squared Error (RMSE):", rmse)

# Calculate R-squared (R2) score
r2 = r2_score(y_test, gb_predictions)
print("R-squared (R2) Score:", r2)

Mean Absolute Error (MAE): 0.2210071896344121
Mean Squared Error (MSE): 0.07875900850305091
Root Mean Squared Error (RMSE): 0.280640354373798
R-squared (R2) Score: 0.9377288383285162

# Calculate Mean Absolute Error (MAE)
mae = mean_absolute_error(y_test, xgb_predictions)
print("Mean Absolute Error (MAE):", mae)

# Calculate Mean Squared Error (MSE)
mse = mean_squared_error(y_test, xgb_predictions)
print("Mean Squared Error (MSE):", mse)

# Calculate Root Mean Squared Error (RMSE)
rmse = mean_squared_error(y_test, xgb_predictions, squared=False)
print("Root Mean Squared Error (RMSE):", rmse)

# Calculate R-squared (R2) score
r2 = r2_score(y_test, xgb_predictions)
print("R-squared (R2) Score:", r2)

Mean Absolute Error (MAE): 0.21938442243495956
Mean Squared Error (MSE): 0.07824639518480786
Root Mean Squared Error (RMSE): 0.2797255712029343
R-squared (R2) Score: 0.9381341383370102

```

